

New Brunswick, Canada · 2026

New Brunswick Responsible AI Framework

A practitioner-built, community-owned governance standard for secure, trustworthy, and accountable AI — built for NB's people, not just its technology.

VERSION	Draft v1.1 — Open Consultation
STATUS	Voluntary Standard — Not Enacted Legislation
LICENCE	Creative Commons Attribution 4.0 International (CC BY 4.0)
ALIGNED TO	EU AI Act · NIST AI RMF 1.0 · ISO 42001 · Bill C-8 · OCAP® · UNDRIP · SANS AISMM (May 2026)
AUTHOR	Joie Catipon — Independent AI Governance Practitioner, New Brunswick, Canada
JURISDICTION	New Brunswick, Canada

VERSION STATUS — OPEN FOR PUBLIC CONSULTATION

This is a practitioner-drafted voluntary standard, not enacted provincial legislation. It operates as a procurement condition, professional baseline, and policy proposal pending provincial legislative action and formal consultation. It will be revised following structured consultation with First Nations communities, labour, civil society, and the NB technology sector. Where AI systems engage federal jurisdiction, this framework complements, and does not supplant, federal obligations.

RELATIONSHIP TO THE EU AI ACT — ADAPTATION, NOT ADOPTION

This framework draws on the EU AI Act's risk-tiered architecture as a starting point — not a wholesale import. The EU AI Act was designed for a single market with pan-European enforcement bodies that Canada does not yet have. Where the EU model is adapted here, it is calibrated to NB's provincial jurisdiction and the operational realities of NB's bilingual, SME-driven, public-service-reliant economy.

§ 1 · PREAMBLE

Why New Brunswick Needs Its Own AI Framework

Artificial intelligence is being deployed across New Brunswick's public and private sectors at an accelerating pace — from government service chatbots and cybersecurity threat detection to sentiment analysis tools used by healthcare and social service providers. Without a clear, jurisdiction-specific governance framework, these deployments risk causing harm to citizens, eroding public trust, and exposing organizations to legal and reputational liability.

This framework provides New Brunswick consultants, technology vendors, public sector bodies, and regulated organizations a practical, principled roadmap for deploying AI systems that are **secure, transparent, trustworthy, fair, reliable, traceable, explainable, and accountable** — anchored in the lived realities of NB's bilingual, geographically diverse, and public-service-reliant population.

SCOPE OF APPLICATION

- Provincial and municipal government services and digital transformation projects
- Cybersecurity operations including AI-driven threat detection and incident response
- AI chatbots and virtual assistants interfacing with NB residents
- Sentiment analysis tools applied to social media, call centre, or public feedback data
- Healthcare, social services, and education AI applications
- IT consultancies deploying AI under CyberSecure NB or related provincial programs

BILINGUALISM REQUIREMENT

All AI systems deployed in public-facing NB contexts must respect the Official Languages Act. AI outputs, explanations, and human oversight mechanisms must be available in both English and French. Sentiment analysis tools must be validated against both language corpora to avoid disparate accuracy across linguistic communities.

FIRST NATIONS PEOPLES: RIGHTS-HOLDERS, NOT JUST PROTECTED CLASSES

New Brunswick is home to fifteen First Nations communities across two nations — Wolastoqiyik (Maliseet) and Mi'kmaq peoples — whose rights predate and exist independently of provincial governance structures. This framework explicitly recognizes that First Nations communities in NB are **rights-holders and governance participants** in AI deployment decisions affecting their peoples, territories, or data — not merely a demographic category to be protected from algorithmic bias.

OCAP® PRINCIPLES — MANDATORY FOR AI SYSTEMS INVOLVING FIRST NATIONS DATA

Any AI system that collects, processes, stores, or generates inferences from data about NB First Nations communities must be governed by the OCAP® principles: Ownership (the community owns its data collectively) · Control (the community controls how data is used) · Access (the community has access to data held about it) · Possession (the community maintains physical custody of data). These are not aspirational — they are the established Indigenous data sovereignty standard in Canadian law and practice.

- **Nation-to-Nation consultation required:** AI systems affecting First Nations communities must be developed in consultation with the relevant Nation(s) prior to procurement — not as a post-deployment notification
- **Free, Prior and Informed Consent (FPIC):** Where AI systems could affect First Nations rights, lands, or governance practices, FPIC consistent with UNDRIP applies
- **No extractive AI:** AI systems must not extract, aggregate, or commercialize knowledge, language, cultural expressions, or community data without explicit consent from the relevant Nation's governance body
- **Indigenous language AI:** Wolastoqey and Mi'kmaq languages are living languages. AI systems that process Indigenous language content must be developed with, not merely for, the relevant language communities
- **Framework co-development commitment:** This framework will be revised following structured Nation-to-Nation consultation. First Nations input will materially shape subsequent versions.

ACKNOWLEDGEMENT OF LIMITATIONS

This v1.1 draft was developed by a GRC practitioner, not in co-design with First Nations communities. That is a limitation that must be named, not obscured. The framework will not be complete until it has been shaped by Wolastoqiyik and Mi'kmaq governance bodies directly. We actively invite that engagement.

BILINGUAL AI — INTERNATIONAL REFERENCE: INDIA'S SUVAS MODEL

Canada's bilingual governance challenge in NB has a useful international parallel in India's SUVAS (Supreme Court Vidhik Anuvaad Software) — an AI-powered machine-assisted translation system launched in 2019 by India's Supreme Court to translate judicial documents between English and eleven Indian vernacular languages.

SUVAS is not directly applicable to NB's English-French bilingual requirement — it was designed for a different language family and legal domain. However, its **governance architecture** is precisely what NB should require for any AI translation used in legal or government contexts:

- **Human-in-the-loop post-editing:** SUVAS generates AI draft translations reviewed and refined by trained human experts (including retired judges and senior law clerks) before publication — no raw AI translation is used for consequential legal text
- **Domain-specific training corpus:** Models are trained exclusively on judicial documents — judgments, orders, statutes, pleadings — not general-purpose text, ensuring domain accuracy
- **Continuous feedback loop:** Corrections made during human post-editing are fed back into the system for incremental accuracy improvement over time
- **Scale and validation:** By 2023, SUVAS had translated over 31,000 of the Supreme Court's 36,000 judgments into scheduled languages, demonstrating the viability of AI-assisted bilingual governance at scale

NB BILINGUAL AI GOVERNANCE — SUVAS LESSON APPLIED

The SUVAS model directly informs NB's bilingual AI requirement: any AI system generating text, decisions, or explanations in English or French for NB residents must follow the SUVAS governance pattern — AI drafts, human domain experts validate, errors improve the model, and no raw AI output is used for consequential administrative or legal text without human review. This applies equally to Acadian French dialect content, which requires validation by Acadian French language experts, not just standard French speakers.

CYBERSECURE NB: THE PROVINCIAL AI & CYBERSECURITY GRANT PROGRAM

The **CyberSecure NB** program, delivered through the **CBDC** network, reimburses eligible SMEs up to **\$3,000** for projects integrating cybersecurity and AI into business operations — the primary vehicle through which this framework's principles reach NB's SME ecosystem.

DETAIL	SPECIFICS
Delivery	Administered by local CBDC offices across New Brunswick
Reimbursement	Up to \$3,000 for eligible professional services and implementation
Eligibility	For-profit entity registered in NB with at least \$30,000 annual revenue
How it works	Engage a service provider · pay · submit invoice + report to CBDC
Commitment	Maintain cybersecurity or AI strategy for at least six months after participation

DEADLINE NOTICE

The application deadline for some streams was previously March 31, 2026. Contact your local CBDC office directly to confirm current availability.

SME RECIPIENT OBLIGATIONS UNDER CYBERSECURE NB

A critical governance gap: the CyberSecure NB program scrutinizes eligible costs but not responsible deployment. When an NB SME deploys an AI system, that business becomes the **operator**. The consultant may have left; the vendor may be offshore. The SME owner makes consequential decisions with AI outputs — and its customers, employees, or community members bear the consequences if it fails.

SME Baseline Obligations

01 Know your risk tier. Confirm in writing which risk tier your AI system falls under before your consultant leaves.

02 Understand what data your AI processes. What data does the system collect? Is it stored in Canada? Who has access?

03 Maintain a human decision-maker. For any AI output affecting a customer or employee — the AI advises; you decide.

04 Keep the vendor honest. Your AI vendor's contract must include: what the system does, what happens when it fails, and how you exit if it causes harm.

05 Tell people. If your business uses AI to screen applications, respond to inquiries, or analyze feedback — disclose it.

<i>AI USE CASE</i>	<i>RISK TIER</i>	<i>SME MINIMUM OBLIGATIONS</i>
AI Chatbot (customer service)	Tier 2	Disclose AI use · offer human escalation · log interactions · monthly output review
Sentiment analysis (feedback)	Tier 2	Inform subjects · do not use as sole basis for decisions · retain audit log
AI-assisted hiring / recruitment	Tier 1	Full human review · bias audit · candidates informed AI was used
Predictive inventory / forecasting	Tier 3	Document the system · validate outputs periodically against actuals
AI bookkeeping / accounting	Tier 3	Human sign-off on financial outputs · data must remain in Canadian jurisdiction

SCOPE BOUNDARIES: WHAT THIS VERSION DOES NOT COVER

<i>OUT-OF-SCOPE AREA</i>	<i>WHY IT MATTERS</i>	<i>PLANNED</i>
Environmental harms	AI energy consumption implicates NB Power's grid and NB's climate commitments	v2.0 — NB Power input required
Labour displacement	Workers displaced by AI-funded automation deserve retraining and notification obligations	v2.0 — NBFL and PSAC engagement
Youth & children's AI safety	Relational dependence, cognitive impact, harmful content in youth-facing AI tools	v2.0 — Education NB input
AI and democratic integrity	Generative AI disinformation and electoral interference in NB political environment	v2.0 — Elections NB engagement

§ 2 · FOUNDATIONAL PRINCIPLES

Nine Core Principles of NB Responsible AI

These principles are the testable standards against which AI systems deployed in NB must be assessed. Two principles — **Indigenous Sovereignty** and **Environmentally Conscious** — were added in v1.1 following initial scrutiny review.

PRINCIPLE	STANDARD	VER
Secure by Design	AI systems must be architected with cybersecurity controls embedded from inception, aligned with CyberSecure NB zero-trust principles and NIST CSF 2.0	Core
Transparent	Users must be informed when interacting with an AI system. Operators must document training data, decision influence, and output generation	Core
Fair & Non-Discriminatory	AI systems must be tested for bias across NB Human Rights Act protected grounds, including language, race, gender, disability, and age — across both official language communities	Core
Reliable & Robust	Systems must perform consistently across edge cases, adversarial inputs, and operational changes. Failure modes must be identified before deployment	Core
Traceable	All AI-influenced decisions must be auditable. Logs of model inputs, outputs, versions, and human overrides must be retained and accessible to oversight bodies	Core
Explainable	Affected individuals and oversight bodies must receive plain-language explanations of AI outputs. XAI requirements are proportional to risk tier	Core
Accountable	A named individual or body must bear legal and operational accountability for each AI system. Responsibility cannot be delegated to a vendor or the AI system itself	Core
Indigenous Sovereignty	First Nations peoples in NB are rights-holders in AI governance. AI systems affecting Wolastoqiyik and Mi'kmaq communities must respect OCAP® principles, FPIC, and UNDRIP — and be developed with, not merely for, those communities	v1.1
Environmentally Conscious	AI deployment decisions must acknowledge energy consumption and carbon implications. Large-scale AI infrastructure choices carry environmental costs that are a legitimate subject of governance. Addressed fully in v2.0	v1.1

§ 3 · RISK CLASSIFICATION

AI Risk Tier Model for New Brunswick

Adapted from the EU AI Act's four-tier risk hierarchy, tailored for NB's public sector, consultancy, and cybersecurity operational contexts. Risk tier determines conformity requirements, documentation obligations, and human oversight thresholds.

TIER	LEVEL	NB EXAMPLES	KEY OBLIGATIONS
Tier 0 — PROHIBITED	Banned Outright	Real-time biometric mass surveillance of NB citizens; social scoring systems; subliminal manipulation	Banned outright. No deployment permitted. Cannot be procured by NB public bodies.
Tier 1 — HIGH RISK	Mandatory Full Oversight	AI in child welfare; parole/sentencing tools; CI protection AI; automated government hiring; healthcare diagnostics	Mandatory conformity assessment, bias audit, human-in-the-loop, registration, incident reporting, DPIA required
Tier 2 — LIMITED RISK	Transparenc y & Logging	GNB portal chatbots; citizen sentiment analysis; AI-assisted cybersecurity triage; procurement advisory AI	Transparency obligations required; logging and audit trail; human review of flagged outputs; annual review cycle
Tier 3 — MINIMAL RISK	Voluntary Standards	AI spam filters; grammar/translation tools; internal knowledge management; code review AI for internal teams	Voluntary code of conduct. Basic documentation recommended. No mandatory conformity assessment.

PRACTITIONER NOTE — WHEN IN DOUBT, CLASSIFY UP

Correct risk tier classification is one of the most consequential decisions in any AI engagement. Deploying a Tier 1 system under Tier 2 obligations — whether through misunderstanding or client pressure — exposes NB residents to unmitigated harm and the practitioner to serious professional and reputational consequences.

S 4 · IMPLEMENTATION PILLARS

Five Pillars of NB AI Implementation

These pillars translate the nine principles into actionable technical and governance requirements. Each pillar has a designated accountability owner within the deploying organization.

1

SECURITY ARCHITECTURE

Cybersecurity-First AI Design

Accountability: CISO or designated AI Security Lead

- Conduct AI-specific threat modelling (prompt injection, model inversion, data poisoning, adversarial inputs) before deployment
- Apply zero-trust access controls to AI model APIs, training pipelines, and inference endpoints
- Implement runtime monitoring for model drift, anomalous output patterns, and unauthorized access attempts
- Maintain a **Model Card** and **System Card** for every deployed model, updated at each version change
- For RAG architectures and AI chatbots: validate retrieval sources, apply output filtering, and log all retrieval-to-response chains
- Perform penetration testing of AI interfaces at minimum annually (quarterly for Tier 1 systems)
- **Data residency:** AI systems processing NB public sector, health, or First Nations data must default to Canadian data residency
- **Foundation model governance:** Deployers using commercial foundation models remain fully responsible for outputs — 'we're just calling an API' does not transfer accountability to the model vendor

2

TRANSPARENCY & DISCLOSURE

Honest AI Interaction Standards

Accountability: Communications or Digital Services Lead

- All chatbot and virtual assistant interfaces must display a visible, plain-language disclosure in both English and French
- Automated decisions that materially affect a person must be identified as AI-generated, with a human review pathway offered
- Sentiment analysis tools must disclose to data subjects that their communications are subject to automated scoring, where legally required
- System documentation must be available to oversight bodies upon request within 30 days
- AI-generated content used in government communications must be disclosed as such

3

FAIRNESS & BIAS MANAGEMENT

Equitable AI for All New Brunswickers

Accountability: Equity Officer, legal counsel, or designated bias auditor

- Conduct pre-deployment bias assessments across all NB Human Rights Act protected grounds including race, national origin, age, disability, sex, and gender identity
- Validate model performance parity across English and French language inputs — especially for sentiment analysis and NLP models
- Test for disparate impact on First Nations communities, rural residents, and seniors
- Maintain a Bias Incident Register; review quarterly; remediate within defined SLAs
- Do not use postal code, language, or community affiliation as proxy variables for protected characteristics

4

EXPLAINABILITY & AUDITABILITY

Traceable AI Decision Chains

Accountability: Data Governance Officer or AI Compliance Lead. Log retention: 5 years (Tier 1), 2 years (Tier 2)

- Implement structured audit logs: input → model version → output → any human override, with timestamps and operator identifiers
- For high-risk decisions: provide human-readable explanation (SHAP, LIME, or equivalent XAI) to affected parties within 10 business days of request
- Maintain version control of all deployed models; document what changed between versions and why
- Conduct post-deployment monitoring: track accuracy, drift, and outcome disparity on a rolling basis
- Establish an AI Incident Reporting process: define what constitutes an AI incident, escalation path, and public disclosure thresholds

5

GOVERNANCE & HUMAN OVERSIGHT

Keeping Humans Accountable

Accountability: Named AI Accountability Owner — designated at Deputy Minister or C-suite level. Cannot be delegated to a vendor, technical team, or the AI system itself

- Designate an **AI Accountability Owner** (named individual) for every deployed system — accountable for compliance, incidents, and remediation
- Establish an AI Review Board with representation from legal, privacy, equity, and technical functions
- Require human-in-the-loop review for all Tier 1 AI outputs before action is taken on a citizen
- Implement a formal AI procurement checklist: vendors must provide conformity documentation, security attestations, and bias testing evidence
- Annual AI portfolio review: assess risk tier reassignments, system retirements, and emerging risks from model updates or scope creep
- Establish whistleblower protection for employees reporting AI-related harms or policy violations in NB public sector bodies
- **SME handover obligation:** When the engagement ends, provide the SME client with a written AI System Summary — risk tier, data processed, what the AI does, what to monitor, who to call if something goes wrong

§ 5 · APPLIED USE CASE GUIDANCE

NB-Specific AI Use Case Guidance

5A · SENTIMENT ANALYSIS IN NB CONTEXTS

RISK TIER

Typically Tier 2 (Limited Risk) for citizen feedback analysis. Escalates to Tier 1 (High Risk) if outputs influence eligibility decisions, law enforcement prioritization, or employment decisions.

- Validate accuracy on NB-specific French and English corpora — Acadian French dialect patterns differ significantly from Parisian French training data
- Test for sarcasm, colloquialisms, and code-switching patterns (Chiac) common in bilingual NB communities
- Establish minimum F1 score thresholds per language and demographic segment before deployment — document parity gaps
- Do not use sentiment scores as sole input to any consequential decision about a person
- Inform data subjects when their communications are sentiment-scored (PIPEDA / NB Privacy Act obligations)
- Establish a human review pathway for outputs scored at confidence thresholds below 0.75
- **Prohibited:** No real-time sentiment analysis of NB residents in public spaces without explicit consent and legal authority
- **Prohibited:** No sentiment scoring of children's communications in education contexts without specific ethical review

5B · AI CHATBOTS IN NB PUBLIC AND GOVERNMENT SERVICES

RISK TIER

Typically Tier 2 (Limited Risk) for information and service navigation chatbots. Escalates to Tier 1 if the chatbot influences eligibility, benefits entitlement, or clinical referral decisions.

- Display bilingual AI disclosure at chatbot launch — not buried in footer terms
- Offer clear escalation path to human agent at any point in the conversation
- Disclose data retention practices: how long chat logs are stored, by whom, for what purpose
- For RAG-based chatbots: validate retrieval sources against authoritative GNB content; implement document freshness controls
- Implement prompt injection defenses — especially critical for public-facing chatbots processing untrusted user input
- Rate-limit and monitor for abuse patterns (data harvesting, jailbreak attempts)
- Do not store personally identifiable information in chatbot session logs beyond operational necessity

§ 5C · BIAS CLASSIFICATION & MITIGATION

Understanding and Addressing AI Bias in NB Deployments

Bias in AI systems manifests at multiple stages of the machine learning pipeline. This section draws on **Amii's Principled AI** framework and **IBM's AI Fairness 360 (AIF360)** to give NB deployers a structured approach to bias identification and mitigation.

THE SIX PIPELINE BIAS TYPES (AMII / SURESH & GUTTAG)

BIAS TYPE	WHERE IT ENTERS	NB EXAMPLE	MITIGATION
Historical	Training data reflecting past societal inequities	Hiring AI trained on historical NB decisions may perpetuate underrepresentation of Acadian, First Nations, or rural candidates	Audit training data provenance; apply reweighting or resampling
Representation	Training data underrepresents certain populations	Sentiment analysis trained on English social media data underperforms on Acadian French text	Supplement with NB French-language corpora; set minimum language-parity thresholds
Measurement	Proxy variables less accurate for some groups	Using postal code as creditworthiness proxy encodes rural NB economic disadvantage	Validate proxy variables for differential accuracy; document all proxy choices
Aggregation	Single model applied to heterogeneous populations	Healthcare triage AI applied uniformly without accounting for distinct First Nations health profiles	Stratify models or use group-specific calibration; sub-group performance analysis
Evaluation	Benchmarks don't represent deployment population	Chatbot evaluated only on standard English queries fails on Chiac (NB code-switching) inputs	Build evaluation datasets representative of NB user language patterns
Deployment	Model used in a context it was not designed for	Sentiment tool validated for customer service applied to social services caseworker interactions	Scope AI systems narrowly; flag any scope expansion as a new deployment requiring fresh assessment

IBM AI FAIRNESS 360 (AIF360) — BIAS MITIGATION PIPELINE

<i>STAGE</i>	<i>WHEN APPLIED</i>	<i>AIF360 TOOLS</i>	<i>NB MINIMUM REQUIREMENT</i>
Pre-processi ng	Before model training — on the dataset	Optimized Preprocessing, Reweighting, Disparate Impact Remover	Required for Tier 1; recommended for Tier 2 NLP and sentiment systems
In-processin g	During model training	Adversarial Debiasing, Meta-Fair Classifier, Learning Fair Representations	Required for Tier 1 systems where training pipeline is within the deployer's control
Post-proces sing	After prediction — on model outputs	Equalized Odds Post-Processing, Reject Option Classification, Calibrated Equalized Odds	Required for ALL Tier 1 and Tier 2 systems — this stage is always within deployer control

NB-SPECIFIC BIAS PRIORITY: LANGUAGE REPRESENTATION

Given NB's unique status as Canada's only officially bilingual province, representation bias and evaluation bias related to French and English language parity are the highest-priority bias risks for any NLP, sentiment analysis, or conversational AI system deployed in NB. No such system should be approved for Tier 1 or Tier 2 deployment without documented evidence of equivalent accuracy across both official languages — including validation against Acadian French dialect patterns, not just standard Parisian French training corpora.

§ 5D · CRITICAL INFRASTRUCTURE

AI in New Brunswick's Critical Infrastructure

New Brunswick is not just a participant in Canada's critical infrastructure ecosystem — it is increasingly a national hub for it. AI in critical infrastructure represents the highest-stakes deployment context in this framework.

LIVE LEGISLATIVE CONTEXT — BILL C-8 / CCSPA (2026)

Bill C-8 passed the House of Commons on March 26, 2026, enacting the Critical Cyber Systems Protection Act. CI operators in NB — including energy, telecommunications, transportation, and finance — now face mandatory cyber security program requirements and incident reporting obligations. AI systems in these sectors must meet both this NB framework and Bill C-8's federal obligations.

NB'S CRITICAL INFRASTRUCTURE AI ECOSYSTEM

INSTITUTION / ASSET	RELEVANCE TO AI IN NB CI
Siemens Cybersecurity Centre for Critical Infrastructure	Fredericton — global OT/ICS security research and deployment
Thales National Digital Excellence Centre (NDEC)	Fredericton — critical infrastructure cyber defence
Canadian Institute for Cybersecurity (CIC) at UNB	Canada's first CI cybersecurity research institute — 50+ researchers
NBCC CI-SOC	Saint John — Canada's first CI Security Operations Centre for IT/OT convergence training
NB Power Smart Grid	AI-enabled grid modernization with Siemens; NERC CIP standards enforced by NB EUB

AI RISK SUB-CLASSIFICATION FOR CI CONTEXTS

SUB-TIER	USE CASE	SECTOR EXAMPLES	KEY OBLIGATIONS
1-A · Critical	Autonomous control of physical systems	Grid switching, water treatment dosing, pipeline pressure management	Human-in-loop mandatory for ALL consequential actions; formal safety case; OT penetration testing; IEC 62443 alignment
1-B · High	AI-driven threat detection	SOC AI, SIEM ML, network behavioural analytics	Model drift monitoring; false positive/negative SLAs; human analyst review before automated response; quarterly red-team testing

SUB-TIER	USE CASE	SECTOR EXAMPLES	KEY OBLIGATIONS
1-B · High	Predictive maintenance AI	Failure prediction, maintenance scheduling for energy/transport/water	Explainability required; human sign-off before acting on AI-generated deferrals
2 · Limited	Administrative AI in CI organizations	Procurement AI, HR screening, financial forecasting	Standard Tier 2 obligations plus elevated supply chain scrutiny

FIVE CI-SPECIFIC AI SECURITY REQUIREMENTS

01 OT Architecture Segregation. AI inference must not occur on the same network segment as live control systems without a validated air-gap or unidirectional gateway. Apply IEC 62443 as the baseline.

02 AI Supply Chain Risk Management. Maintain a complete AI Bill of Materials (AI-BOM). No AI component from an unvetted vendor may be deployed without a formal security assessment. AI supply chain incidents are reportable under Bill C-8.

03 Adversarial Attack Preparedness. Test AI systems against prompt injection, model inversion, membership inference, data poisoning, evasion attacks, and model extraction. Red-team exercises: annually for Tier 1-B, semi-annually for Tier 1-A.

04 AI Failure Mode & Continuity Planning. Every AI system in CI must have a documented and tested manual override procedure. Operational plans cannot assume AI availability during a cyber incident.

05 AI Incident Reporting. AI-related CI incidents must be reported to CCCS via My Cyber Portal within Bill C-8 timeframes. Also notify CyberSecure NB and, for energy, the NB EUB. Post-incident reviews within 30 days.

§ 6 · GOVERNANCE STRUCTURE

NB AI Governance Architecture

BODY	MANDATE
NB AI Policy & Standards Office (Proposed — GNB)	Develop and maintain the NB AI Framework; publish mandatory AI procurement standards; maintain a public NB AI Systems Registry for Tier 1 and Tier 2 systems; issue guidance on emerging AI risks including generative AI and agentic systems
Office of the Information & Privacy Commissioner of NB (Expanded Mandate)	Review AI-related DPIAs for Tier 1 systems; receive and investigate complaints from NB residents regarding AI decisions; coordinate with OPC on cross-jurisdictional deployments. Note: binding AI-specific enforcement powers require legislative expansion of the IPC's mandate — this framework recommends that expansion.
CyberSecure NB (Technical Authority)	Publish and maintain AI-specific cybersecurity technical standards; provide threat intelligence on emerging AI attack vectors; conduct or certify security assessments of Tier 1 AI systems; support incident response when AI systems are involved in cybersecurity events

CONSULTANT PRACTICE STANDARDS

WRITTEN BY A PRACTITIONER, FOR PRACTITIONERS

This framework was drafted by an NB GRC and AI Trust practitioner who works in the same consultancy ecosystem it describes. The following are not punitive obligations — they are the professional baseline that any responsible AI consultant in NB should already be applying, offered as a shared standard for what clients should expect from any AI engagement in NB.

- Complete a risk tier classification for every AI system proposed or deployed — documented and discussed with the client before work begins
- Provide clients with a plain-language AI System Summary covering: purpose, risk tier, data inputs, key limitations, and escalation contacts — in both official languages where the client serves bilingual populations
- Evaluate proposed AI vendors against NB framework requirements and document findings in the SOW or contract schedule
- Notify the client promptly upon discovery of any AI security incident, model failure, or evidence of bias that affects NB residents or their data
- Do not recommend, deploy, or white-label AI systems that cannot provide training data provenance, security documentation, or bias testing results — if a vendor cannot answer these questions, that is itself a risk finding
- Where a client engagement involves First Nations communities or data, apply OCAP® principles and recommend Nation-to-Nation consultation before any AI system is scoped, not after it is built
- **SME handover obligation:** When the engagement ends, provide the SME client with a written AI System Summary they can actually use after you leave

S 7 · IMPLEMENTATION ROADMAP

Phased Adoption Timeline

A realistic, sequenced implementation plan. The roadmap reflects the current federal policy environment — no binding Canadian AI statute is in force as of mid-2026 — and is calibrated toward voluntary adoption and stakeholder consultation as the near-term levers.

Summer-Fall 2026
Current Phase

Open Consultation & Community Feedback

Publish framework publicly. Invite feedback from NB tech sector, GNB staff, privacy advocates, First Nations governance bodies, labour organizations, and civil society. Initiate Nation-to-Nation outreach to Wolastoqiyik and Mi'kmaq nations. Seek academic peer review from UNB, Université de Moncton, and Mount Allison.

Fall 2026 Policy
Submission

Formal Submission — Revised Draft v2.0

Submit revised framework incorporating consultation feedback to GNB, NBIF, and federal AI strategy consultation channels as a model for sub-national AI governance in bilingual Canadian jurisdictions. Publish final bilingual version in English and French.

Q1 2027 Voluntary
Adoption

Toolkits & Pilot Programs

Release self-assessment toolkit, risk classification guide, and Model Card template. Pilot with 3–5 GNB departments. Engage NBCC and UNB for AI ethics curriculum alignment.

Q2-Q3 2027 Mandatory
Standards

GNB Procurement Standards

AI procurement checklist becomes a condition of GNB technology RFPs. CyberSecure NB publishes AI security baseline. IPC begins accepting AI-related DPIA filings. First Nations consultation findings incorporated into v2.0.

Q4 2027 Public
Registry

NB AI Systems Registry Launch

Public registry of Tier1 and Tier2 AI systems deployed by GNB goes live. Annual transparency report published. Incident reporting mechanism activated.

2028 Biennial Review

First Framework Review — Scope Expansion

Review against federal AI strategy implementation and EU AI Act enforcement lessons. Expand scope to cover environmental harms, labour displacement, and youth AI safety — areas explicitly deferred from v1.1.

§ 8 · REFERENCES & FRAMEWORK ALIGNMENT

Framework Alignment Map

SOURCE	KEY ELEMENTS REFERENCED
EU AI Act (2024/1689) — International	Risk tier architecture (Articles 5, 6, 50) — adapted, not adopted wholesale; Transparency (Article 13); Human oversight mandate (Article 14)
NIST AI RMF 1.0 — USA / NIST	GOVERN, MAP, MEASURE, MANAGE functions; AI risk identification; Trustworthiness characteristics taxonomy. Explicit pillar mapping in §8b below.
ISO/IEC 42001:2023 — International Standard	AI management system requirements; Organizational accountability controls; AI impact assessment methodology
Federal AI Strategy (2026) — Canadian Federal	AIDA / Bill C-27 died on order paper January 2025 — successor expected; ISED National AI Sprint Report (February 2026); Evan Solomon appointed Canada's first Minister for AI and Digital Innovation (May 2025)
Bill C-8 / CCSPA (2026) — Canadian Federal	Critical Cyber Systems Protection Act — passed House of Commons March 26, 2026; Mandatory cyber security programs for designated CI operators; CCCS Critical Infrastructure Resilience initiative (April 2026)
Directive on Automated Decision-Making — Federal Treasury Board	Algorithmic impact assessment tool; Decision automation levels 1–4; Notice, explanation, and appeal rights; Peer review and audit requirements
OCAP® Principles & UNDRIP — Indigenous Rights & Data	OCAP® — First Nations Information Governance Centre (Ownership, Control, Access, Possession); UNDRIP Act (Canada, 2021); FPIC
India SUVAS — Bilingual AI Governance Reference	Human-in-the-loop post-editing model for AI translation; domain-specific training corpus approach; continuous feedback loop for accuracy improvement — governance architecture applicable to NB bilingual AI requirements
New Brunswick Legislation — Provincial	Protection of Personal Information Act; NB Human Rights Act; Official Languages Act; NB Energy and Utilities Board Act; CyberSecure NB / CBDC Program
Amii Principled AI & IBM AIF360 — Bias & Fairness	Suresh & Gutttag six pipeline bias types; IBM AI Fairness 360 — 70+ fairness metrics and 10 mitigation algorithms; Amii Datasheets for Datasets, Model Cards, and Algorithmic Impact Assessments
SANS AI Security Maturity Model™ (May 2026)	Three pillars: Protect AI, Utilize AI, Govern AI. Five maturity stages. Agentic AI, Non-Human Identity, and the Principle of Least Agency. Explicit NB stage mapping in §8c below.

§ 8B · NIST AI RMF 1.0 — EXPLICIT PILLAR MAPPING

The NIST AI Risk Management Framework organizes AI risk management into four core functions: **GOVERN**, **MAP**, **MEASURE**, and **MANAGE**. The following maps each of this framework's five pillars to the NIST AI RMF function it primarily addresses.

NB PILLAR	PRIMARY NIST FUNCTION	SECONDARY	KEY NIST ALIGNMENT
Pillar 1 — Cybersecurity-First AI Design	MANAGE	MAP	MANAGE 1.0 (risks prioritized and addressed); MAP 1.0 (context established, risks identified); threat modelling and penetration testing align to MANAGE 2.0 risk treatment actions
Pillar 2 — Honest AI Interaction Standards	GOVERN	MANAGE	GOVERN 1.0 (policies for trustworthy AI); GOVERN 4.0 (commitment to transparency); disclosure requirements align to NIST Trustworthiness characteristic: Transparency
Pillar 3 — Fairness & Bias Management	MEASURE	MAP	MEASURE 2.0 (AI risks analyzed, including bias and fairness); MAP 5.0 (impacts on individuals, communities); bias incident register aligns to MEASURE 4.0 (risk measurement results documented)
Pillar 4 — Explainability & Auditability	MEASURE	GOVERN	MEASURE 1.0 (AI risk measurement approaches identified); GOVERN 6.0 (policies for transparency and accountability); audit logs and XAI align to NIST Trustworthiness: Explainability and Accountability
Pillar 5 — Governance & Human Oversight	GOVERN	MANAGE	GOVERN 1.0–6.0 (full governance function); AI Accountability Owner maps to GOVERN 2.0; human-in-the-loop requirements align to MANAGE 3.0

NIST AI RMF FUNCTION DEFINITIONS

GOVERN: Cultivates organizational practices that make AI risk management a priority — policies, culture, roles, accountability structures.

MAP: Categorizes AI risks in context — AI system identification, stakeholder engagement, risk identification across the AI lifecycle.

MEASURE: Analyzes and assesses AI risks quantitatively and qualitatively — bias evaluation, performance monitoring, risk measurement documentation.

MANAGE: Prioritizes and addresses AI risks — risk treatment, incident response, and residual risk monitoring.

§ 8C · SANS AI SECURITY MATURITY MODEL™ — NB ALIGNMENT

The **SANS AI Security Maturity Model™** (released May 12, 2026) provides a five-stage, three-pillar operational roadmap from ad hoc AI use to fully governed and secured AI programs. Where the SANS model tells NB organizations how mature their security posture is, this framework tells them what governance obligations apply at each tier.

The Three SANS Pillars — Mapped to NB Framework

<i>SANS PILLAR</i>	<i>SANS DEFINITION</i>	<i>NB FRAMEWORK ALIGNMENT</i>
Protect AI	Securing AI systems from external and internal threats — model integrity, data pipeline security, inference endpoint protection, adversarial attack resilience	Pillar 1 (Cybersecurity-First AI Design); §5d CI requirements 1–4; threat modelling, penetration testing, OT architecture segregation, AI-BOM
Utilize AI	Governing how AI tools are adopted — shadow AI, approved model registries, data governance for AI inputs, responsible use policies	Pillar 2 (Transparent Disclosure); SME Recipient Obligations (§1); foundation model governance requirement; data residency requirements
Govern AI	Organizational accountability, risk management, compliance, and oversight — policies, risk registers, audit trails, incident response, regulatory reporting	Pillar 5 (Governance & Human Oversight); Pillar 4 (Explainability & Auditability); AI Accountability Owner; AI Review Board; Bill C-8 incident reporting

The Five SANS Maturity Stages — NB Context and Obligations

<i>STAGE</i>	<i>SANS DEFINITION</i>	<i>NB ORGANIZATION PROFILE</i>	<i>MINIMUM NB OBLIGATIONS</i>
Stage 1 Unaware / Ad Hoc	No formal AI security policy. AI tools adopted without visibility or governance. Sensitive data moving into public models without controls.	Most NB SMEs receiving CyberSecure NB grants. Organizations deploying first AI chatbot or sentiment tool without a governance framework.	Complete SME Baseline Obligations (§1). Conduct risk tier classification immediately. Apply Tier 3 minimum obligations. Do not proceed to Tier 1 or Tier 2 deployments without uplift.
Stage 2 Aware / Reactive	AI use is visible but governance is reactive. Some policies exist but not enforced consistently.	NB mid-sized organizations with an IT function but no dedicated AI governance role. GNB departments with informal AI use but no AI procurement standard.	Implement Pillar 1 and Pillar 2. Assign an AI Accountability Owner. Establish a Bias Incident Register. Apply Tier 2 obligations to all chatbot and sentiment deployments.

STAGE	SANS DEFINITION	NB ORGANIZATION PROFILE	MINIMUM NB OBLIGATIONS
Stage 3 Defined / Managed	Formal AI security program in place. Risk assessments conducted. Controls documented and measured. Red teaming at least quarterly.	NB critical infrastructure operators (energy, health, finance). GNB departments with formal digital governance. Organizations under Bill C-8 obligations.	All five pillars active. Full Tier 1 conformity assessment complete. AI-BOM maintained. Bias audits documented. Audit logs retained per Pillar 4. Incident reporting operational.
Stage 4 Optimized	Continuous AI risk monitoring. Adversarial simulation in production. Regulatory evidence generated automatically. Agent-to-agent interactions logged and audited.	NB CI operators targeting full Bill C-8 compliance and EU AI Act audit readiness. Organizations with agentic AI deployments.	Apply Principle of Least Agency to all agentic AI systems: each agent has a documented identity, scoped permissions per task, and session-level audit trail. NHI governance required for all AI agents.
Stage 5 Optimizing / Adaptive	AI security posture updates in near real-time. Security controls tested continuously via adversarial simulation. Organization contributes to industry standards.	Aspirational for NB CI hub organizations (Siemens, Thales, CIC/UNB). NB organizations providing AI security services to other sectors.	All NB framework obligations met. Contribute findings to NB AI Systems Registry and proposed NB AI Policy Office. Share threat intelligence with CyberSecure NB and CCCS.

AGENTIC AI – THE NB GOVERNANCE GAP TO WATCH

The SANS AI Security Maturity Model identifies agentic AI and Non-Human Identity (NHI) as the area where most organizations have a real blind spot. The Principle of Least Agency — the agentic counterpart to least privilege — requires that every AI agent in production be granted only the minimum permissions required for the specific task, with a documented identity, scoped authorization, and a session-level audit trail. NB organizations deploying agentic AI systems (AI that takes autonomous actions, calls external APIs, or orchestrates other AI systems) must apply this principle as a non-negotiable security control. This will be incorporated as a formal requirement in v2.0.

COMMUNITY PURPOSE STATEMENT

This framework exists for one reason: AI is being deployed in New Brunswick right now, and the people of this province — residents, workers, First Nations communities, francophones, rural communities, and small businesses — deserve a governance standard that protects them. It is offered freely as a community resource, not a proprietary product. Adapt it, critique it, improve it, and share it.

LICENSE — CREATIVE COMMONS ATTRIBUTION 4.0 INTERNATIONAL (CC BY 4.0)

Copyright (c) 2026 Joie Catipon — Independent AI Governance Practitioner, New Brunswick, Canada.

You are free to share (copy and redistribute in any medium or format) and adapt (remix, transform, and build upon) this material for any purpose — including by governments, organizations, or other communities — provided you give appropriate credit to the original author, provide a link to the licence, and indicate if changes were made. You may not apply legal terms that restrict others from what the licence permits.

Full licence text: creativecommons.org/licenses/by/4.0

New Brunswick Responsible AI Framework · Draft v1.1 · Open for Public Consultation · 2026
Voluntary Standard — Not Enacted Legislation · Subject to revision following First Nations consultation,
civil society review, and community feedback